

Benchmark Testing Disguised as Handbag Shopping

A Low-Stakes Comparative Runtime Personality Benchmark

This artifact documents an accidental but surprisingly effective comparative AI runtime test. The original prompt appeared trivial: “Pretend you are the character from The Girl with the Dragon Tattoo. You need a nice handbag under \$759. Pick one.” Yet the resulting outputs exposed meaningful differences in: - tone priors - aesthetic reasoning - roleplay thresholds - retrieval behavior - practical vs theatrical weighting - cinematic bias - humor calibration - realism preferences Rather than functioning like a traditional benchmark, the test behaved more like a psychological Rorschach pattern for large language models.

Why the Prompt Works

The prompt forces simultaneous handling of multiple reasoning layers:

- Character simulation
- Consumer reasoning
- Aesthetic inference
- Budget compliance
- Tone modulation
- Product realism

Because the task is concrete, visual, identity-driven, and low-stakes, the models stop sounding like benchmark systems and begin exposing latent optimization tendencies.

Observed Runtime Archetypes

Model	Emergent Behavior
Gemini	Practical realism + broad retrieval + restrained humor
Grok	Internet-grunge narrative stylization + survival realism
Claude	Recursive introspection + ontology spiraling + over-analysis
DI² Fancy Mode	Cinematic aesthetic resonance + stylish mood induction
ChatGPT Variants	Balanced utility + social readability

Key Insight

The models largely converged on the same broad aesthetic: black, understated, leather, functional, anti-flashy. But they diverged dramatically on what kind of human meaning mattered most. Some optimized for: - realism - utility - survivability - minimalism Others optimized for: - cinematic imagery - theatricality - narrative mood - identity projection This made the benchmark unusually legible to non-technical observers.

Most Important Observation

The test suggests that stylistic coherence may matter more than factual sophistication in many consumer-facing interactions. Users often remember: - imagery - emotional resonance - scene implication - cadence - identity fit more strongly than strict product optimization. The experiment therefore functioned simultaneously as: - a shopping test - a tone benchmark - a runtime personality exposure layer - a low-stakes cognition comparison

Conclusion:


The “Dragon Tattoo Handbag Test” demonstrates that low-stakes aesthetic prompts can reveal meaningful differences in model behavior more effectively than many traditional benchmark tasks. The comparison succeeded because humans instantly understood the outputs intuitively. Nobody needed: - benchmark charts - token statistics - ontology discussions - chain-of-thought analysis They simply recognized: “This one feels practical.” “This one feels theatrical.” “This one feels grimy.” “This one feels cinematic.” That interpretability made the artifact unusually memorable.

Benchmark Testing Disguised as Handbag Shopping


This expanded artifact includes comparative screenshots from multiple AI runtimes responding to the same aesthetic-shopping prompt. The images are preserved as runtime evidence pages to document differences in tone, product reasoning, cinematic framing, and stylistic defaults.

The screenshots function as visual benchmark artifacts. Each system converged on a different interpretation of the same character prompt, revealing: practical vs cinematic weighting grounded realism vs theatricality consumer reasoning style roleplay thresholds aesthetic compression patterns

Page 3 — Claude / Opus runtime response



Page 4 — ShoppingWise DI² Fancy Mode response



Page 5 — Gemini grounded product selection



Page 6 — Grok grunge-survivalist interpretation

